

Motivation		
Problem <i>What problem are we solving? For example, a tedious & repetitive manual task?</i> <i>insufficient information to make a decision?</i> <i>Who are the target users and beneficiaries?</i> <i>Why is the current solution insufficient?</i> <i>How could the user experience be improved?</i>	Business Value <i>How will this reduce costs (e.g., speed, fewer errors)?</i> <i>How will this increase revenue (e.g., new product feature, retention)?</i> <i>What is the scale of impact (number of users, volume)?</i> <i>Does this enable future projects (e.g., reusing data infrastructure)?</i> <i>Does this create a competitive advantage or is needed to just keep up?</i>	Success Criteria <i>Definition of done and measurable KPIs.</i> <i>Guardrail metrics that must not degrade.</i> <i>Current baselines for comparison.</i> <i>How does ML success (e.g., high model accuracy) translate into business success (e.g., revenue)?</i>

Approach		
Deliverables <i>Software system to make predictions continuously or insights from a one-off analysis?</i>	Integration <i>How does the solution fit into existing workflows?</i> <i>What input data is required and where does it come from?</i> <i>What outputs are produced and how are they utilized?</i>	Build or Buy? <i>Does this require unique domain knowledge?</i>

Required Resources & Biggest Risks		
Data <i>How much data is available (incl. rare events & labels)?</i> <i>Do subject matter experts believe the data is sufficient, or are key variables missing?</i> <i>How easy is it to access and consolidate the data?</i> <i>What effort is required to build a pipeline for continuous data collection?</i>	Deployment & Maintenance <i>Where should the model run (cloud, on-prem, edge)?</i> <i>What infrastructure or hardware (e.g., GPUs) is needed?</i> <i>How will the model be monitored, retrained, and maintained over time?</i> <i>Will the solution scale as data and usage grow?</i>	Risks & Governance <i>What regulations apply (e.g., GDPR, HIPAA, EU AI Act)?</i> <i>Are there risks around data privacy, fairness, or transparency?</i> <i>What are the consequences if the model makes an error?</i> <i>Should humans remain in the loop for critical decisions?</i> <i>Could users try to game or exploit the system?</i> <i>What barriers might prevent user trust or adoption?</i> <i>Under what conditions should the project be stopped (kill switch criteria)?</i>
ML <i>Is there a simpler solution without ML?</i> <i>Has a similar problem been solved with ML before?</i> <i>Can a human solve the task with the given input data?</i>	People <i>Who owns the solution and its outcomes?</i> <i>What expertise is required, and do we need support from other teams or external partners?</i> <i>Will users need training to adopt it effectively?</i>	