

Machine Learning Course Workbook

Introduction

ML is everywhere!

Where (else) do you use ML in your everyday life incl. work?

-

ML history: Why now?

What accelerated the rise of ML in the last few years?

-

-

-

What is the difference between ANI and AGI?

What is ML?

Define ML:

What do you need to create an ML-powered product (i.e., value)?

-

-

-

-

AI and ML Researchers, Statisticians, and Data Scientists all use a certain set of tools.

What is the difference between...

- ML vs. AI?
- ML vs. Deep Learning?
- ML vs. Statistics?
- ML vs. Data Science?

How do machines “learn”?

Describe the different learning strategies:

- Unsupervised Learning:
- Supervised Learning:
- Reinforcement Learning:

When should you use ML?

When should you not use ML?

For which kinds of problems does ML have a high chance of success and when is the outcome uncertain?

What distinguishes an ML project from a data science project (in terms of deliverables)?

In what ways can you create value with ML?

Solving problems with ML: Workflow

What are the 3 main steps to create value with ML?

- 1.
- 2.
- 3.

What are the two main deployment possibilities for an ML model and when should you use which?

-
-

Which tasks take up most of a Data Scientist's time?

Data & Preprocessing

What are "features" and what are "labels"?

- Features:
- Labels:

What does structured and unstructured data look like?

- Structured Data:
- Unstructured Data:

Garbage in, garbage out!

What do you think are the most common ways in which datasets in your organization are messy?

Which concrete next steps could your organization take to improve their data quality?

ML Solutions: Overview

What does the output of the different algorithm categories look like for one data point?

- Dimensionality Reduction:
- Anomaly Detection:
- Clustering:
- Classification:
- Regression:

What are the benefits of breaking down a complex input-output problem into simpler subproblems?

Supervised Learning Basics

What is the difference between a regression and a classification problem?

What is the stupid baseline you should always compare a regression model against?

What is the stupid baseline you should always compare a classification model against?

Avoiding Common Pitfalls

When should you absolutely evaluate a model with the balanced accuracy metric?

What does it mean for a model to over- or underfit?

Why might a model still be wrong, even if it generates correct predictions for new data points?

What are “Adversarial Attacks”?

Why can it happen that a model discriminates?

What is the difference between data and concept drift?

Conclusion

According to Andrew Ng, what are the 5 steps for a successful AI transformation of a company?

- 1.
- 2.
- 3.
- 4.
- 5.

Where do you think your organization stands in this AI transformation process?