

Machine Learning Course Workbook

Introduction

ML is everywhere!

Where (else) do you use ML in your everyday life incl. work?

-

ML history: Why now?

What accelerated the rise of ML in the last few years?

-

-

-

What is the difference between ANI and AGI?

What is ML?

Define ML:

What do you need to create an ML-powered product (i.e., value)?

-

-

-

-

AI and ML Researchers, Statisticians, and Data Scientists all use a certain set of tools.

What is the difference between...

- ML vs. AI?
- ML vs. Deep Learning?
- ML vs. Statistics?
- ML vs. Data Science?

How do machines “learn”?

Describe the different learning strategies and what their requirements (in terms of data) are:

- Unsupervised Learning:
- Supervised Learning:
- Reinforcement Learning:

When should you use ML?

In what ways can ML create value?

When should you not use ML?

Which kind of ML problems have a high chance of success and when is the outcome uncertain?

Solving problems with ML: Workflow

What are the 3 main steps to create value with ML?

- 1.
- 2.
- 3.

What should you check before starting an ML project?

-
-
-
-

What are the two deployment options for an ML model and when should you use which?

-
-

Which tasks take up most of a Data Scientist's time?

Data & Preprocessing

What are "features" and what are "labels"?

- Features:
- Labels:

What does structured and unstructured data look like?

- Structured Data:
- Unstructured Data:

Garbage in, garbage out!

What do you think are the most common ways in which datasets in your organization are messy?

Which concrete next steps could your organization take to improve their data quality?

ML Solutions: Overview

What does the output of the different algorithm categories look like for one data point?

- Dimensionality Reduction:
- Anomaly Detection:
- Clustering:
- Regression:
- Classification:
- Recommender Systems/Information Retrieval:

What are the benefits of breaking down a complex input-output problem into simpler subproblems?

Avoiding Common Pitfalls

With which stupid baseline should you compare regression and classification models respectively?

When is it a really bad idea to evaluate a classification model with the accuracy metric?

What does it mean for a model to over- or underfit?

Why can a model still be wrong, even though it generates correct predictions for data points from the testset?

What are “Adversarial Attacks”?

In what ways can a biased model negatively affect users?

How can you check whether a model discriminates?

What is the difference between data and concept drift?

What could be reasons for data or concept drift in your domain / next project?

Conclusion

According to Andrew Ng, what are the 5 steps for a successful AI transformation of a company?

- 1.
- 2.
- 3.
- 4.
- 5.

Where do you think your organization stands in this AI transformation process?